


# Cell nuclei segmentation: support vector machine

*Final presentation by:  
Michelle Emmert, Juan Hamdan, Laura Sanchis, und Gloria Timm*



**Our results**




**Evaluation**



**Summary and Outlook**



**Our algorithm**



**Data**



# The Dataset

28 images of nuclei :

- ***N2DH-GOWT1***


GFP transfected GOWT1 mouse embryonic stem cells

- ***N2DL-HeLa***


Histone 2B (H2B)-GFP expressing HeLa cells

- ***NIH3T3***


mouse embryonic fibroblast – CD tagged (EGFP)



***N2DH-GOWT1***



***N2DL-HeLa***

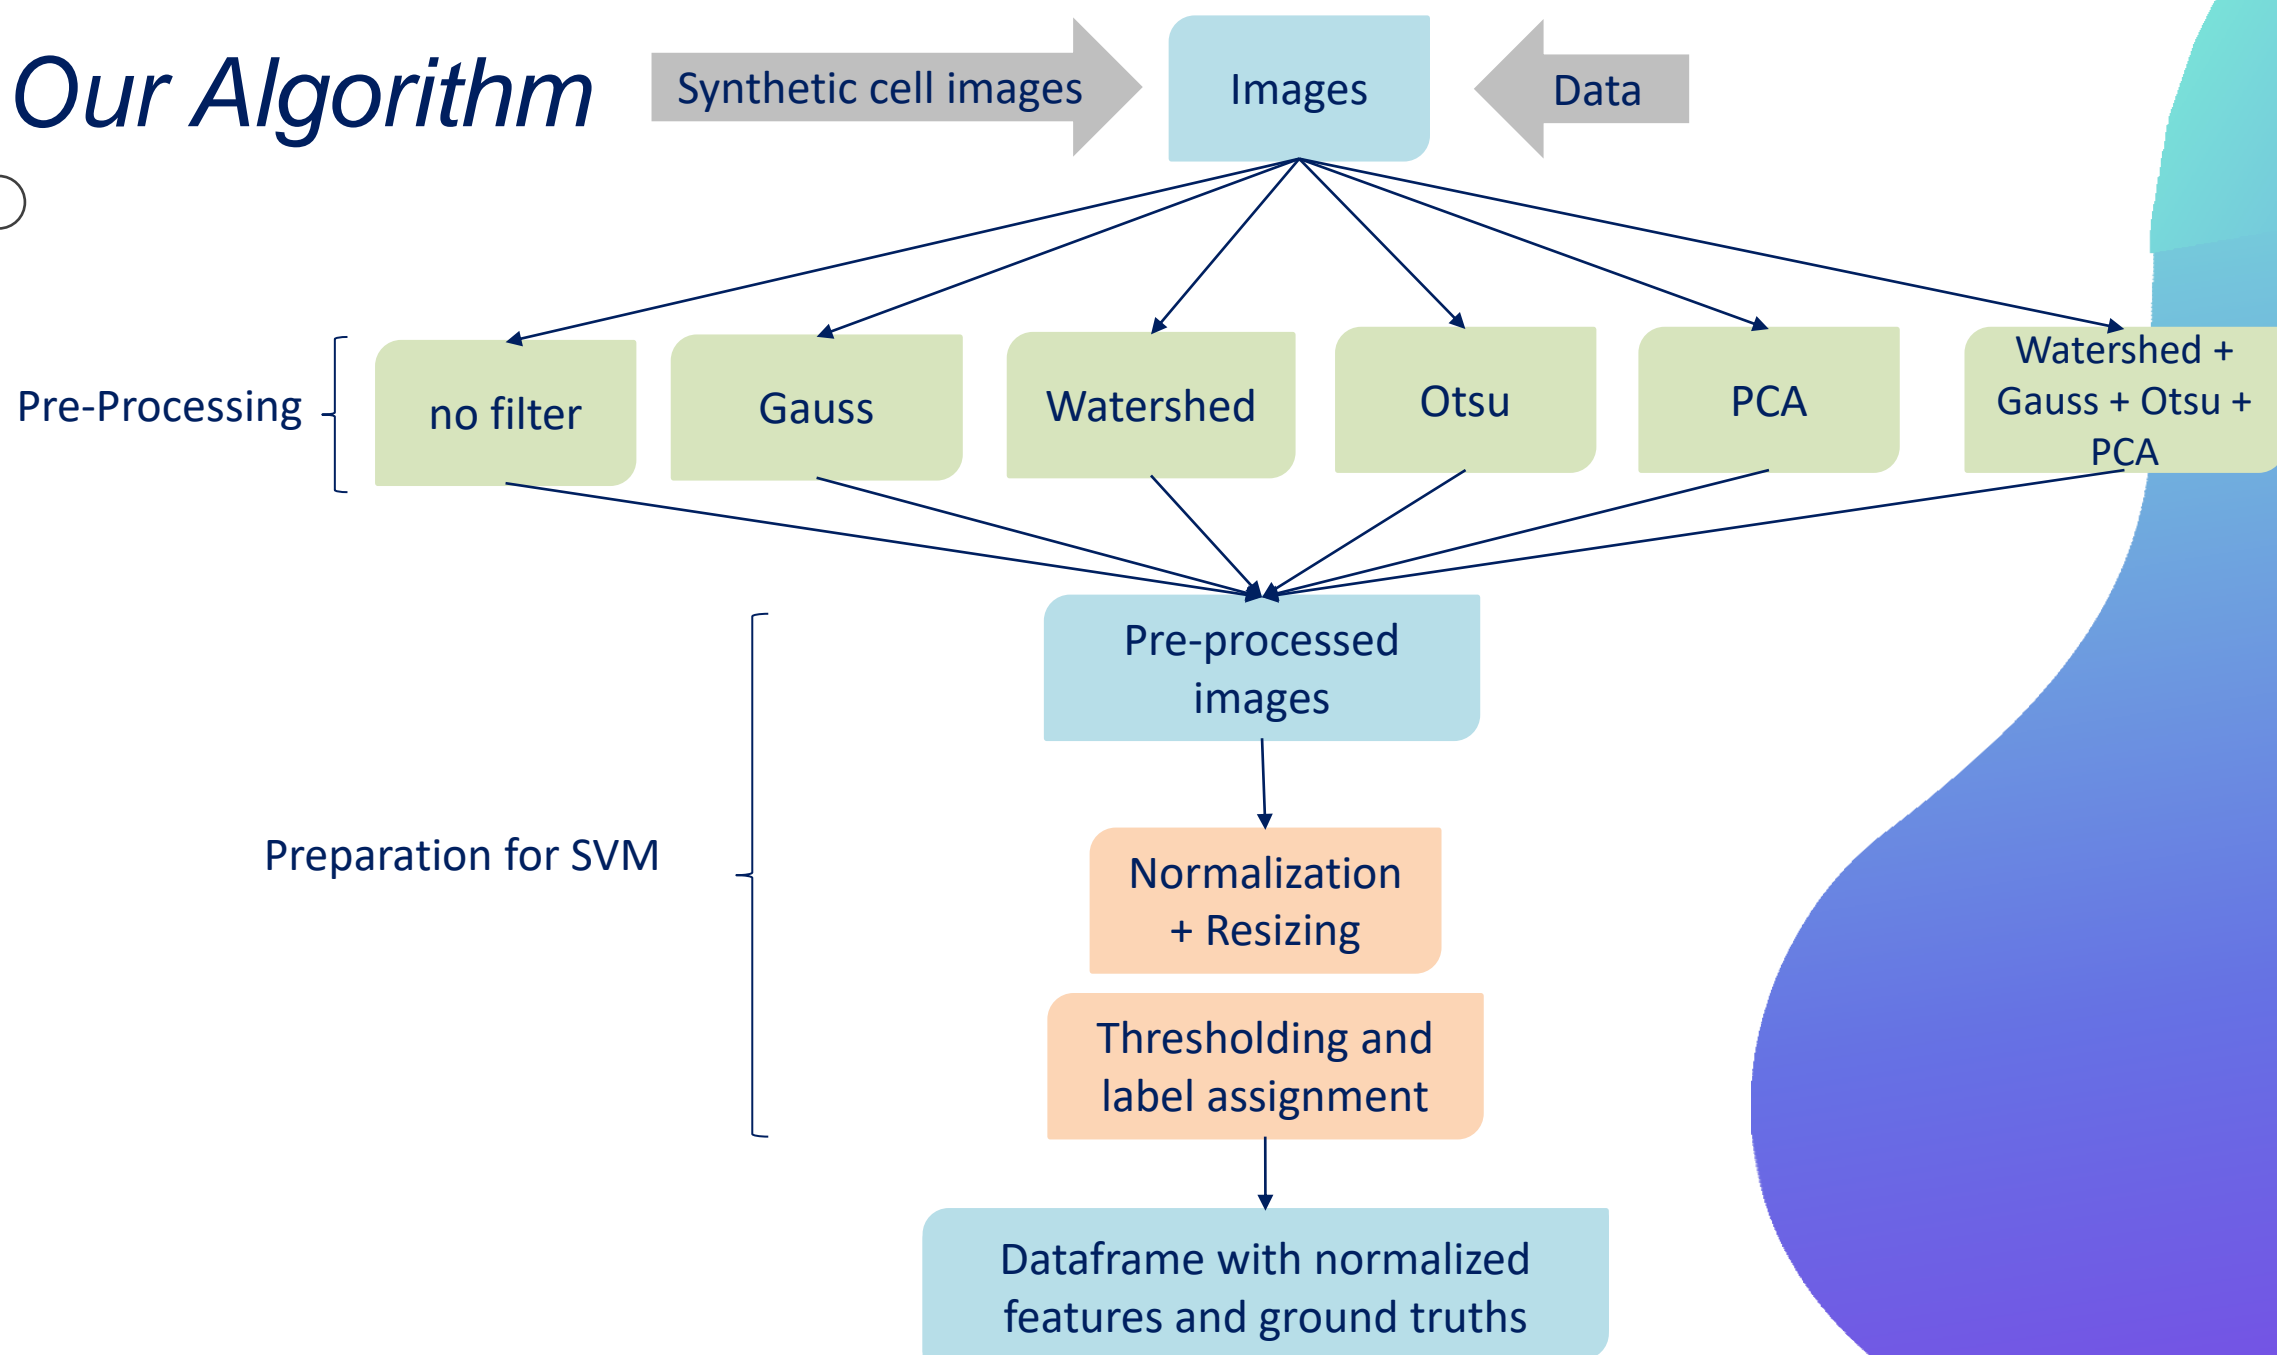


***NIH3T3***

(1) Osuna, E. et al. 2007. Large-Scale Automated Analysis of Location Patterns in Randomly Tagged 3T3Cells

(2) Maska, M. et al. 2014. A benchmark for comparison of cell tracking algorithms

# Our Algorithm



# *Our Algorithm*

Segmentation

Evaluation

Dataframe with normalized  
features and ground truths

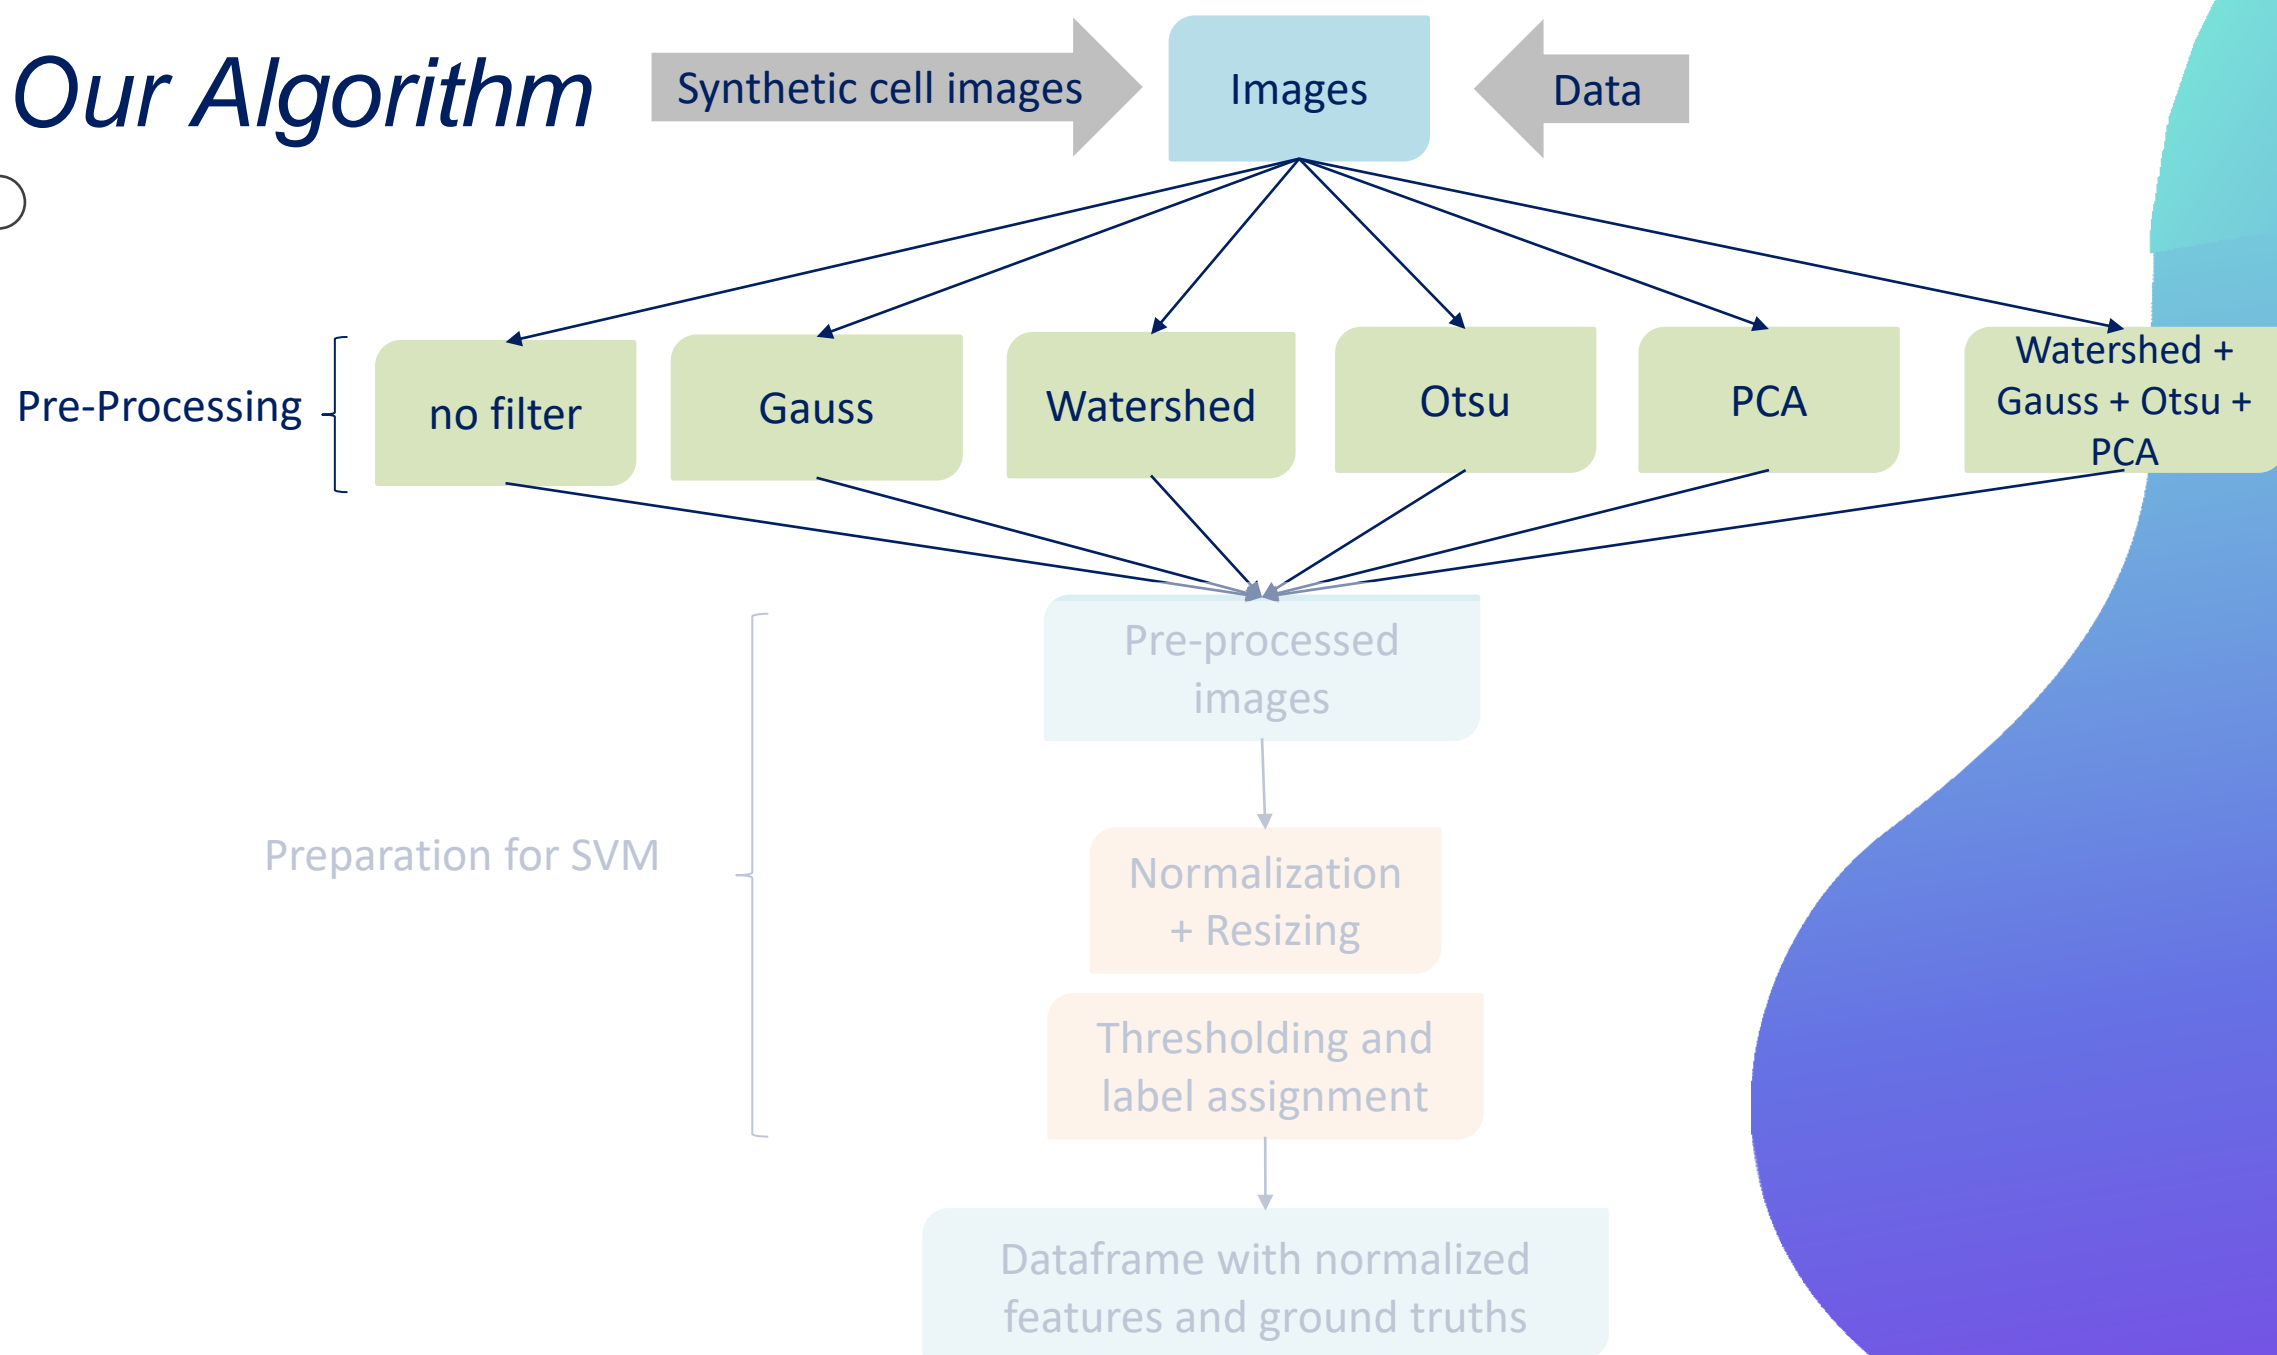
SVM

Segmented images

Dice score



# Our Algorithm







# *Synthetic images*

Goals:

- test the Dice score code
- enrich and enlarge our training data set for the SVM

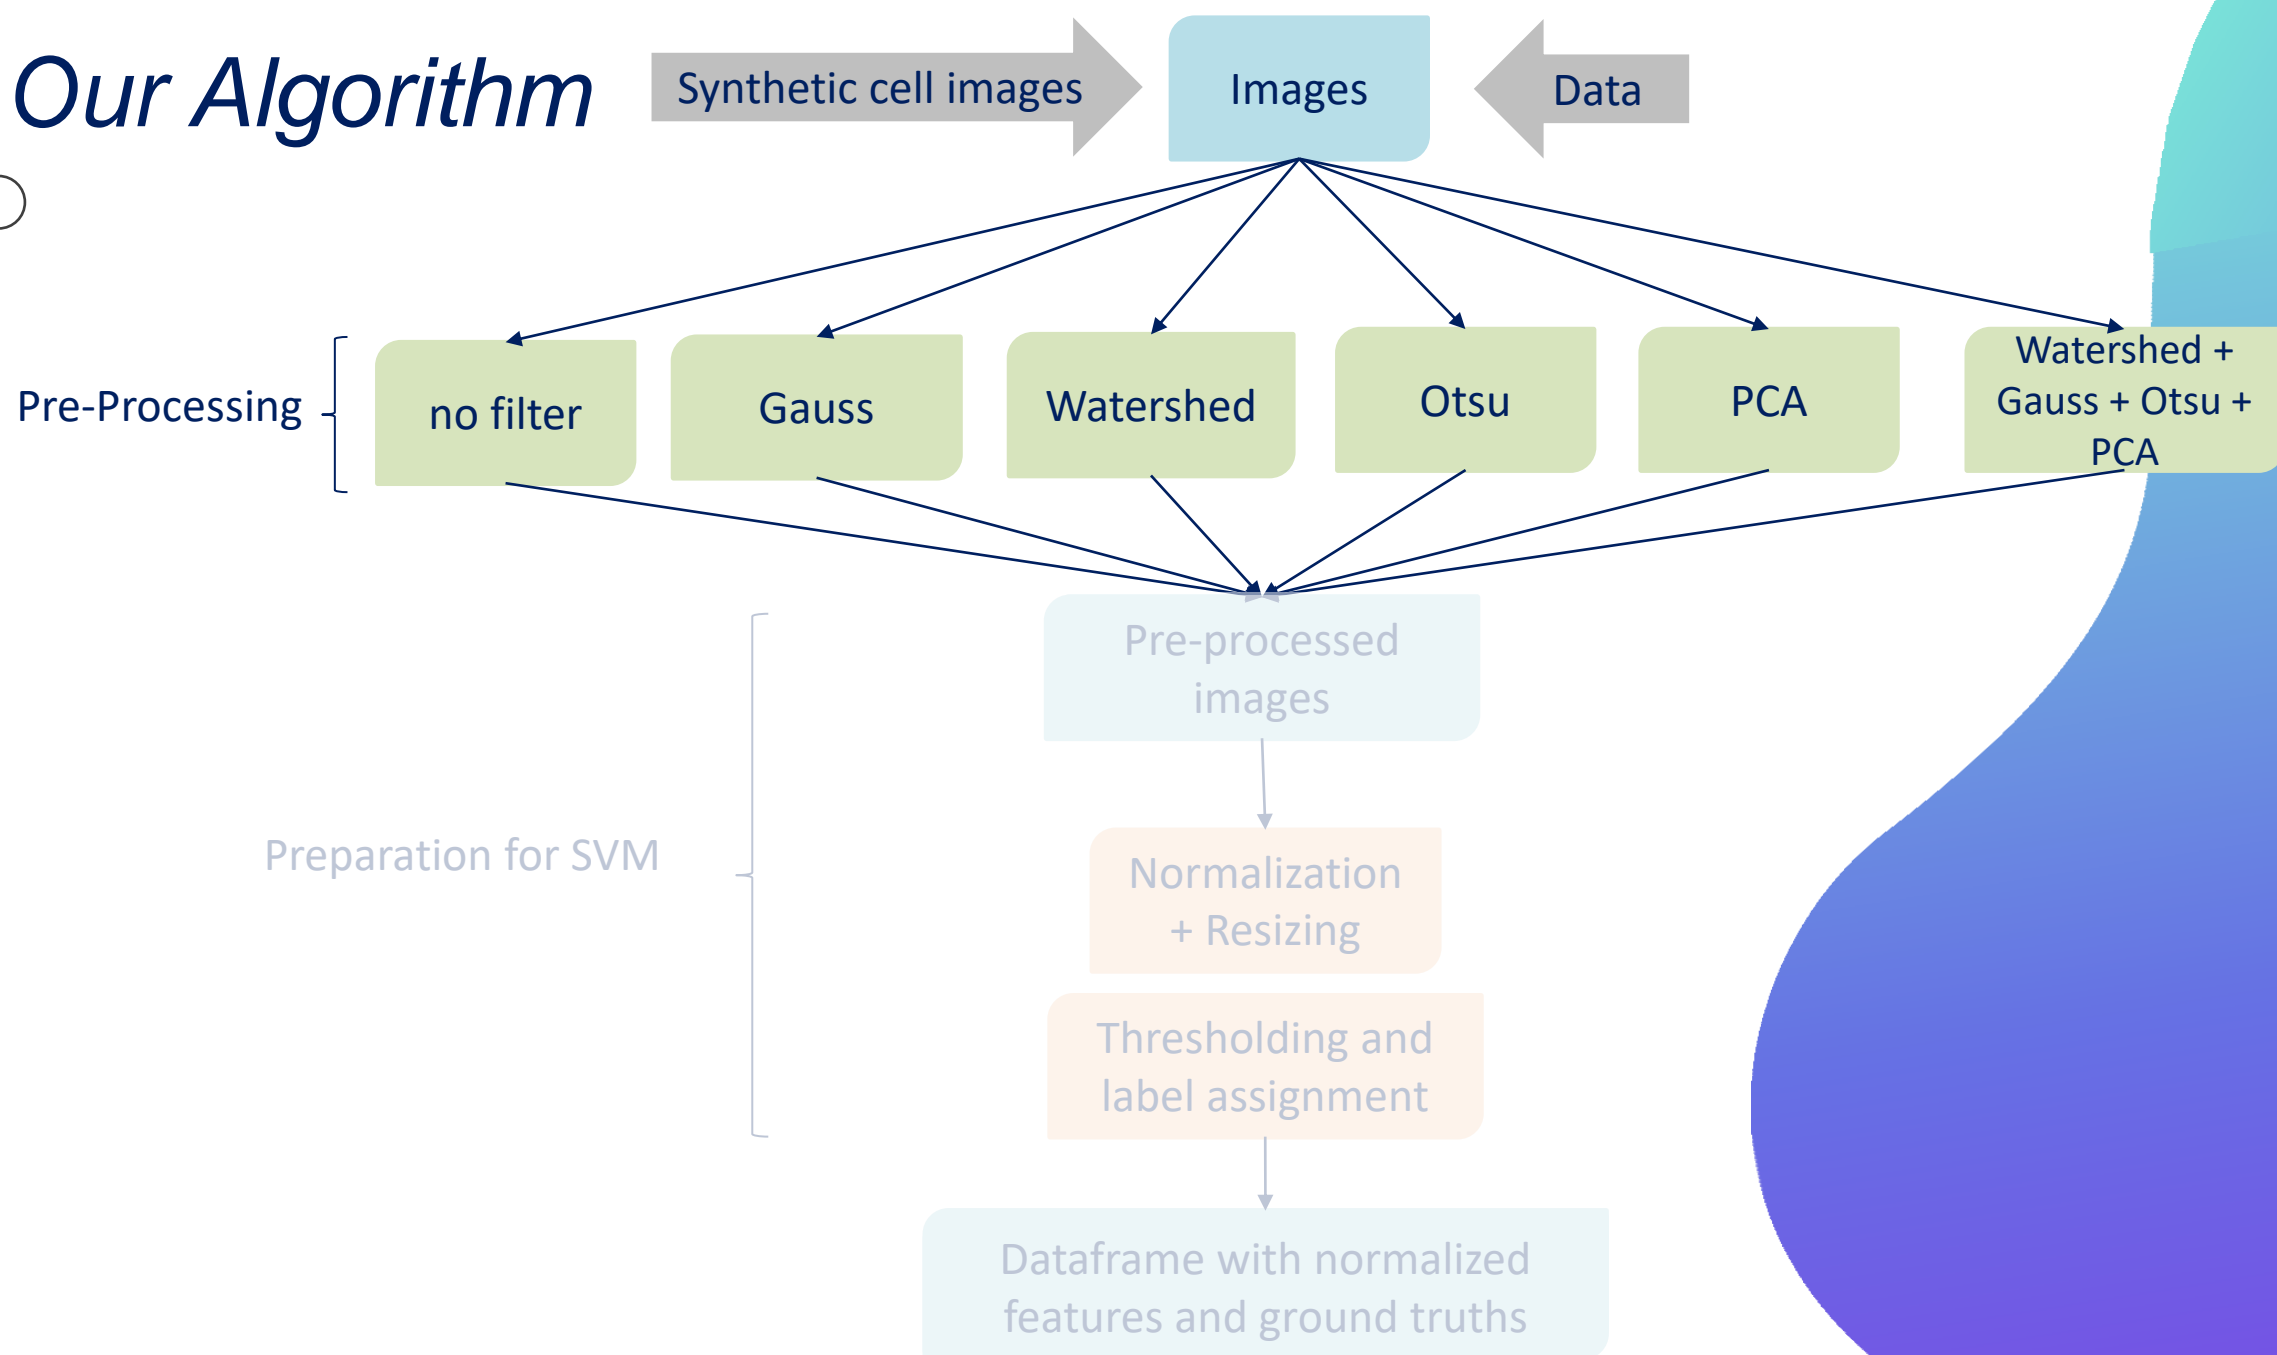
Result:

- did not improve Dice score
- advantage: being able to segment almost all original images





# Our Algorithm





# Pre-processing Methods


Goal:

improve Dice Score of segmentation method through better image quality


Desired effect:

1. average local pixel intensity values
2. separate nuclei which appear fused
3. increase contrast
4. enhance edges


through noise reduction, super-pixel segmentation , thresholding and data selection.



# Pre-processing Methods




Gaussian Filtering



Desired effect:  
noise reduction



# Pre-processing Methods



Watershed Filtering




Desired effect:  
edge enhancement  
separate fused nuclei




<http://www.cmm.mines-paristech.fr/~beucher/wtshed.html>

# Pre-processing Methods




Otsu




Desired effect:

increase contrast and enhance edges

# Pre-processing Methods



PCA




Desired effect:

reduce noise and increase contrast

not used to for dimension reduction!

# Our Algorithm



# *Our Algorithm*

Segmentation

Dataframe with normalized  
features and ground truths

SVM

Segmented images


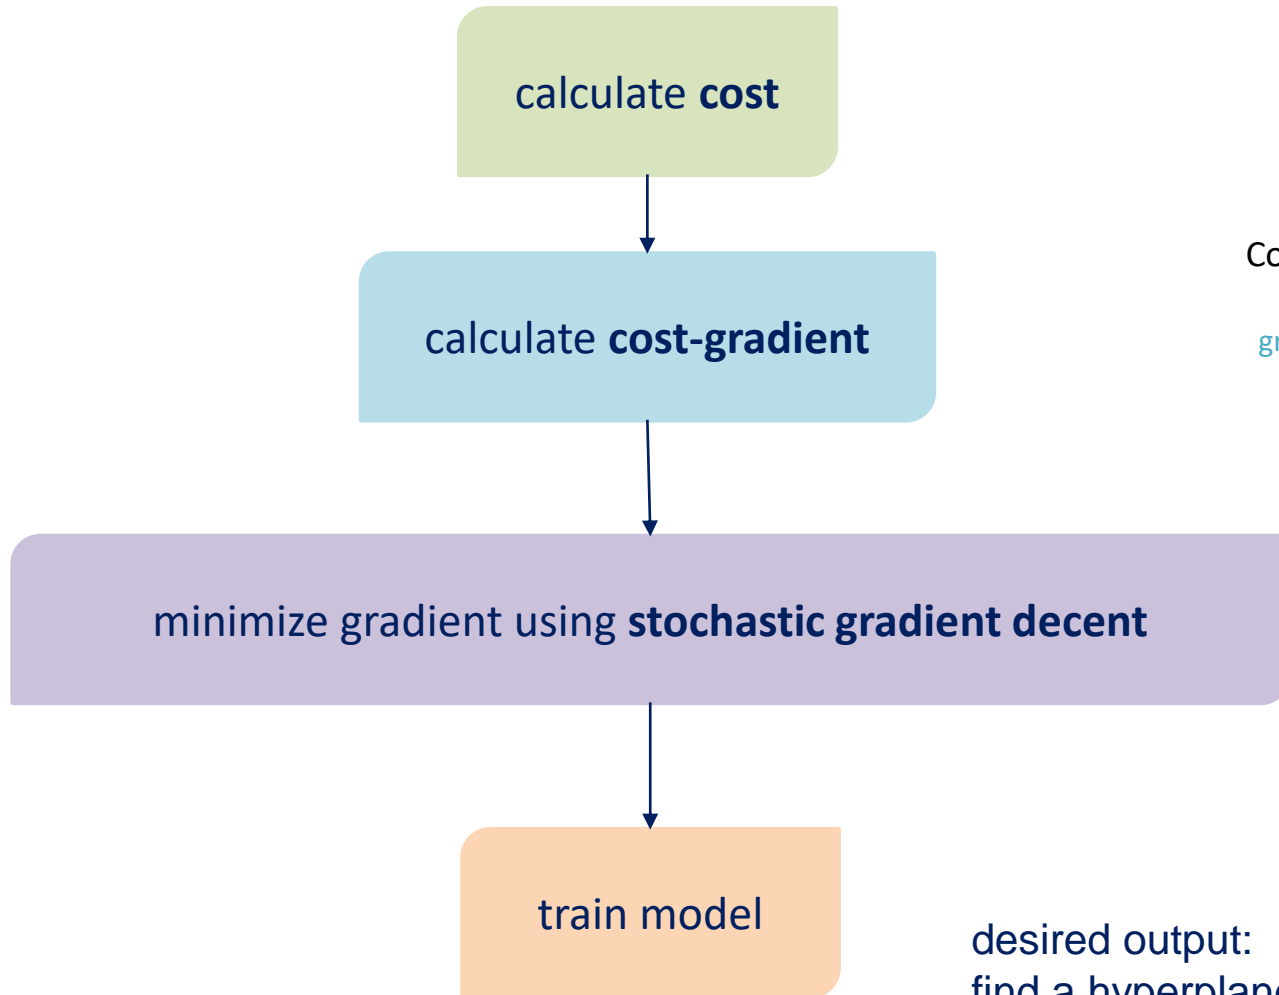
Evaluation

Dice score





# Support vector machine



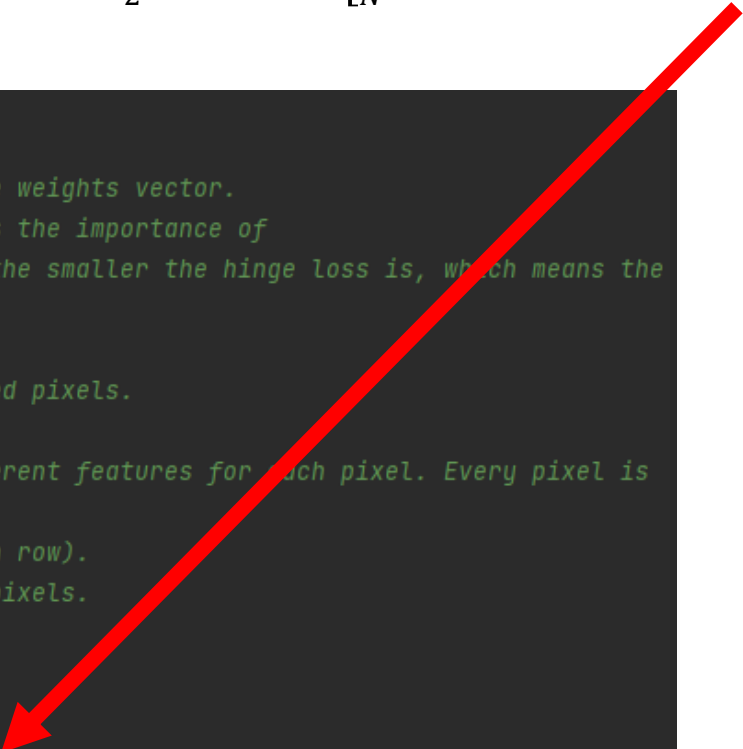
desired output:  
find a hyperplane that separates samples belonging to two classes  
with the largest margin possible, while keeping the misclassification low  
→ minimize loss function

# Support vector machine

calculate cost

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

```
17 def compute_cost(weights, features, labels, soft_margin_factor):
18     """
19     This function calculates the cost for a given feature, label and the weights vector.
20     First, the hinge loss is calculated, a loss function that determines the importance of
21     misclassified pixels. The smaller the soft margin factor (smf) is, the smaller the hinge loss is, which means the
22     SVM margin will be softer.
23     The cost function depends on the weights vector and the hinge loss.
24     :param soft_margin_factor: determines the importance of misclassified pixels.
25     :param weights: Weights vector
26     :param features: Feature vector, has different columns for the different features for each pixel. Every pixel is
27     a row.
28     :param labels: Labels vector, is either 1 or -1 for each pixel (each row).
29     :return: The cost of the current weights vector at classifying the pixels.
30     """
31     # calculate hinge loss
32     number_pixels = features.shape[0]
33     distances_to_hyperplane = 1 - labels * (np.dot(features, weights))
34     distances_to_hyperplane = np.maximum(0, distances_to_hyperplane)
35     hinge_loss = soft_margin_factor * (np.sum(distances_to_hyperplane) / number_pixels)
36     # calculate cost
37     cost = 1 / 2 * np.dot(weights, weights) + hinge_loss
38     return cost
```



# Support vector machine

calculate cost

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

```
17 def compute_cost(weights, features, labels, soft_margin_factor):
18     """
19     This function calculates the cost for a given feature, label and the weights vector.
20     First, the hinge loss is calculated, a loss function that determines the importance of
21     misclassified pixels. The smaller the soft margin factor (smf) is, the smaller the hinge loss is, which means the
22     SVM margin will be softer.
23     The cost function depends on the weights vector and the hinge loss.
24     :param soft_margin_factor: determines the importance of misclassified pixels.
25     :param weights: Weights vector
26     :param features: Feature vector, has different columns for the different features for each pixel. Every pixel is
27     a row.
28     :param labels: Labels vector, is either 1 or -1 for each pixel (each row).
29     :return: The cost of the current weights vector at classifying the pixels.
30     """
31     # calculate hinge loss
32     number_pixels = features.shape[0]
33     distances_to_hyperplane = 1 - labels * (np.dot(features, weights))
34     distances_to_hyperplane = np.maximum(0, distances_to_hyperplane)
35     hinge_loss = soft_margin_factor * (np.sum(distances_to_hyperplane) / number_pixels)
36     # calculate cost
37     cost = 1 / 2 * np.dot(weights, weights) + hinge_loss
38     return cost
```

# Support vector machine

calculate cost

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

```
17 def compute_cost(weights, features, labels, soft_margin_factor):
18     """
19     This function calculates the cost for a given feature, label and the weights vector.
20     First, the hinge loss is calculated, a loss function that determines the importance of
21     misclassified pixels. The smaller the soft margin factor (smf) is, the smaller the hinge loss is, which means the
22     SVM margin will be softer.
23     The cost function depends on the weights vector and the hinge loss.
24     :param soft_margin_factor: determines the importance of misclassified pixels.
25     :param weights: Weights vector
26     :param features: Feature vector, has different columns for the different features for each pixel. Every pixel is
27     a row.
28     :param labels: Labels vector, is either 1 or -1 for each pixel (each row).
29     :return: The cost of the current weights vector at classifying the pixels.
30     """
31     # calculate hinge loss
32     number_pixels = features.shape[0]
33     distances_to_hyperplane = 1 - labels * (np.dot(features, weights))
34     distances_to_hyperplane = np.maximum(0, distances_to_hyperplane)
35     hinge_loss = soft_margin_factor * (np.sum(distances_to_hyperplane) / number_pixels)
36     # calculate cost
37     cost = 1 / 2 * np.dot(weights, weights) + hinge_loss
38     return cost
```

# Support vector machine

calculate cost

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

```
17 def compute_cost(weights, features, labels, soft_margin_factor):
18     """
19     This function calculates the cost for a given feature, label and the weight vector.
20     First, the hinge loss is calculated, a loss function that determines the importance of
21     misclassified pixels. The smaller the soft margin factor (smf) is, the smaller the hinge loss is, which means the
22     SVM margin will be softer.
23     The cost function depends on the weights vector and the hinge loss.
24     :param soft_margin_factor: determines the importance of misclassified pixels.
25     :param weights: Weights vector
26     :param features: Feature vector, has different columns for the different features for each pixel. Every pixel is
27     a row.
28     :param labels: Labels vector, is either 1 or -1 for each pixel (each row).
29     :return: The cost of the current weights vector at classifying the pixels.
30     """
31     # calculate hinge loss
32     number_pixels = features.shape[0]
33     distances_to_hyperplane = 1 - labels * (np.dot(features, weights))
34     distances_to_hyperplane = np.maximum(0, distances_to_hyperplane)
35     hinge_loss = soft_margin_factor * (np.sum(distances_to_hyperplane) / number_pixels)
36     # calculate cost
37     cost = 1 / 2 * np.dot(weights, weights) + hinge_loss
38     return cost
```

# Support vector machine

calculate **cost**



calculate **cost-gradient**

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

$$\nabla_w \text{cost}(w) = \frac{1}{N} \sum_i^n \begin{cases} w & \text{if } \max(0, 1 - y_i * (w * x_i)) = 0 \\ w - C y_i x_i & \text{otherwise} \end{cases}$$

# Support vector machine

calculate cost-gradient

$$\nabla_w \text{cost}(w) = \frac{1}{N} \sum_i^n \begin{cases} w & \text{if } \max(0, 1 - y_i * (w * x_i)) = 0 \\ w - C y_i x_i & \text{otherwise} \end{cases}$$

```
41 def calculate_cost_gradient(weights, features, labels, soft_margin_factor):
42     """
43     This function calculates the gradient of the cost function. It computes its derivative, in order to know in which
44     direction the gradient descent has to go to find the minimum of the cost function.
45     :param soft_margin_factor: determines the importance of misclassified pixels.
46     :param weights: weights vector.
47     :param features: a vector with all our pixels as rows.
48     :param labels: label of the pixels, either +1 or -1.
49     :return: gradient of the cost function with that weight vector.
50     """
51     # In order to iterate correctly, we turn our vectors into arrays.
52     labels = np.array([labels])
53     features = np.array([features])
54
55     # Calculate distance to hyperplane, to classify the pixels. This dot product can be changed depending on the
56     # chosen kernel.
57     distance_to_hyperplane = 1 - (labels * np.dot(features, weights))
58
59     # Create an empty gradient vector, to fill with the gradient of the current pixel.
60     gradient_cost = np.zeros(len(weights))
61
62     for index_pixel, distance_pixel in enumerate(distance_to_hyperplane):
63         # For correctly classified pixels, the current weight vector is maintained
64         if max(0, distance_pixel) == 0:
65             gradient_pixel = weights
66         # For incorrectly classified pixels, the weight vector is corrected in the direction contrary to the gradient.
67         else:
68             gradient_pixel = weights - (soft_margin_factor * labels[index_pixel] * features[index_pixel])
69         gradient_cost += gradient_pixel
70
71     return gradient_cost
```



# Support vector machine

calculate **cost-gradient**

$$\nabla_w \text{cost}(w) = \frac{1}{N} \sum_i^n \begin{cases} w & \text{if } \max(0, 1 - y_i * (w * x_i)) = 0 \\ w - C y_i x_i & \text{otherwise} \end{cases}$$

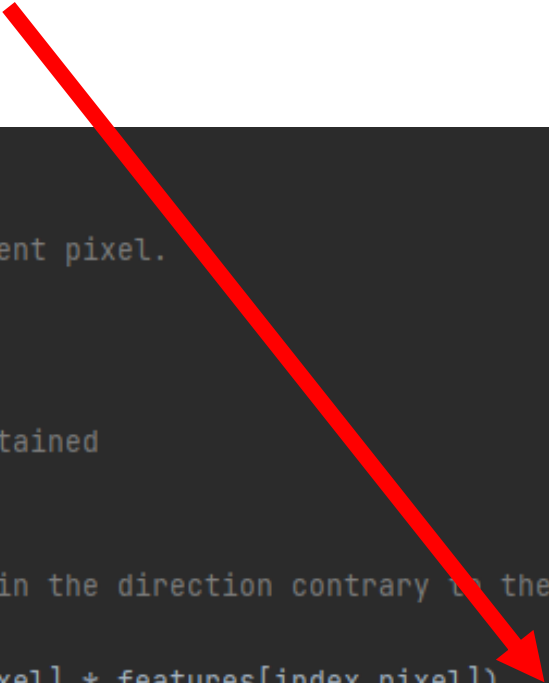
```
41 def calculate_cost_gradient(weights, features, labels, soft_margin_factor):
42     """
43     This function calculates the gradient of the cost function. To compute the gradient, we need to know the distance of each pixel to the hyperplane.
44     """
45     distance_to_hyperplane = 1 - (labels * np.dot(features, weights))
46
47     # Create an empty gradient vector, to fill with the gradient of the current pixel.
48     gradient_cost = np.zeros(len(weights))
49
50     for index_pixel, distance_pixel in enumerate(distance_to_hyperplane):
51         # For correctly classified pixels, the current weight vector is maintained
52         if max(0, distance_pixel) == 0:
53             gradient_pixel = weights
54         # For incorrectly classified pixels, the weight vector is corrected in the direction contrary to the gradient.
55         else:
56             gradient_pixel = weights - (soft_margin_factor * labels[index_pixel] * features[index_pixel])
57         gradient_cost += gradient_pixel
58
59     return gradient_cost
60
61 else:
62     gradient_pixel = weights - (soft_margin_factor * labels[index_pixel] * features[index_pixel])
63     gradient_cost += gradient_pixel
64
65 return gradient_cost
```

# Support vector machine

calculate **cost-gradient**

$$\nabla_w \text{cost}(w) = \frac{1}{N} \sum_i^n \begin{cases} w & \text{if } \max(0, 1 - y_i * (w * x_i)) = 0 \\ w - C y_i x_i & \text{otherwise} \end{cases}$$

```
41 def calculate_cost_gradient(weights, features, labels, soft_margin_factor):
42     """
43     This function calculates the gradient of the cost function. To compute the gradient, we need to know the distance
44     of each pixel to the hyperplane.
45     """
46     distance_to_hyperplane = 1 - (labels * np.dot(features, weights))
47
48     # Create an empty gradient vector, to fill with the gradient of the current pixel.
49     gradient_cost = np.zeros(len(weights))
50
51     for index_pixel, distance_pixel in enumerate(distance_to_hyperplane):
52         # For correctly classified pixels, the current weight vector is maintained
53         if max(0, distance_pixel) == 0:
54             gradient_pixel = weights
55         # For incorrectly classified pixels, the weight vector is corrected in the direction contrary to the gradient.
56         else:
57             gradient_pixel = weights - (soft_margin_factor * labels[index_pixel] * features[index_pixel])
58             gradient_cost += gradient_pixel
59
60     return gradient_cost
61
62 else:
63     gradient_pixel = weights - (soft_margin_factor * labels[index_pixel] * features[index_pixel])
64     gradient_cost += gradient_pixel
65
66 return gradient_cost
```



# Support vector machine

calculate **cost**

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

calculate **cost-gradient**

$$\nabla_w \text{cost}(w) = \frac{1}{N} \sum_i^n \begin{cases} w & \text{if } \max(0, 1 - y_i * (w * x_i)) = 0 \\ w - C y_i x_i & \text{otherwise} \end{cases}$$

minimize gradient using **stochastic gradient decent**

# Support vector machine

minimize gradient using **stochastic gradient decent**

```
74 def sgd(features, labels, soft_margin_factor, learning_rate, max_epochs):
75     """
76     This function calculates the stochastic gradient descent to minimize our cost function.
77     :param features: all pixels as rows, with n columns, which stand for n features.
78     :param labels: all pixels as rows, only one column per pixel, either +1 or -1.
79     :param soft_margin_factor:
80     :param learning_rate:
81     :param max_epochs: maximum number of cycles to try to find the minimum of the cost function.
82     :return: the weight vector found at the minimum of the cost function, and the history of how the costs have
83             evolved during the calculations.
84     """
85     # Create an empty weight vector that is the same size as the number of columns (features) of a single pixel.
86     weights = np.zeros(features.shape[1])
87     # Define the first cost for our function and an empty history cost list.
88     prev_cost = float("inf")
89     history_cost = []
90     patience = 0
91
92     # Stochastic gradient descent
93     for epoch in range(0, max_epochs):
94         # shuffle to prevent repeating update cycles
95         features_shuffled, labels_shuffled = shuffle(features, labels)
96         for pixel_index, pixel_value in enumerate(features_shuffled):
97             gradient = calculate_cost_gradient(weights, pixel_value, labels_shuffled[pixel_index], soft_margin_factor)
98             weights = weights - (learning_rate * gradient)
99
100         # Calculate cost to evaluate the advance.
101         cost = compute_cost(weights, features, labels, soft_margin_factor)
102         history_cost.append(cost)
103         if epoch % 20 == 0 or epoch == max_epochs - 1:
104             print("Epoch is: {} and Cost is: {}".format(epoch, cost))
105
106         # Stoppage criterion
107         if prev_cost < cost:
108             if patience == 10:
109                 return weights, history_cost
110             else:
111                 patience += 1
112         else:
113             patience = 0
114             prev_cost = cost
115
116     return weights, history_cost
```

# Support vector machine

minimize gradient using **stochastic gradient decent**

```
74 def sgd(features, labels, soft_margin_factor, learning_rate, max_epochs):
75     """
76     This function calculates the stochastic gradient descent to minimize our cost function.
77     for pixel_index, pixel_value in enumerate(features_shuffled):
78         gradient = calculate_cost_gradient(weights, pixel_value, labels_shuffled[pixel_index], soft_margin_factor)
79         weights = weights - (learning_rate * gradient)
80
81     # Calculate cost to evaluate the advance.
82     cost = compute_cost(weights, features, labels, soft_margin_factor)
83     history_cost.append(cost)
84     if epoch % 20 == 0 or epoch == max_epochs - 1:
85         print("Epoch is: {} and Cost is: {}".format(epoch, cost))
86
87     # Stoppage criterion
88     if prev_cost < cost:
89         if patience == 10:
90             return weights, history_cost
91         else:
92             patience += 1
93     else:
94         patience = 0
95     prev_cost = cost
96
97     return weights, history_cost
98
99     patience = 0
100     prev_cost = cost
101
102     return weights, history_cost
```

# Support vector machine


minimize gradient using **stochastic gradient decent**

```
74 def sgd(features, labels, soft_margin_factor, learning_rate, max_epochs):
75     """
76     This function calculates the stochastic gradient descent to minimize our cost function.
77     for pixel_index, pixel_value in enumerate(features_shuffled):
78         gradient = calculate_cost_gradient(weights, pixel_value, labels_shuffled[pixel_index], soft_margin_factor)
79         weights = weights - (learning_rate * gradient)
80
81
82
83
84     # Calculate cost to evaluate the advance.
85     cost = compute_cost(weights, features, labels, soft_margin_factor) ←
86     history_cost.append(cost)
87
88     if epoch % 20 == 0 or epoch == max_epochs - 1:
89         print("Epoch is: {} and Cost is: {}".format(epoch, cost))
90
91
92
93
94     # Stoppage criterion
95     if prev_cost < cost:
96         if patience == 10:
97             return weights, history_cost
98         else:
99             patience += 1
100
101     else:
102         patience = 0
103         prev_cost = cost
104
105
106     return weights, history_cost
107
108     patience = 0
109     prev_cost = cost
110
111     return weights, history_cost
```

# Support vector machine

minimize gradient using **stochastic gradient decent**

```
74 def sgd(features, labels, soft_margin_factor, learning_rate, max_epochs):
75     """
76     This function calculates the stochastic gradient descent to minimize our cost function.
77     for pixel_index, pixel_value in enumerate(features_shuffled):
78         gradient = calculate_cost_gradient(weights, pixel_value, labels_shuffled[pixel_index], soft_margin_factor)
79         weights = weights - (learning_rate * gradient)
80
81
82
83
84     # Calculate cost to evaluate the advance.
85     cost = compute_cost(weights, features, labels, soft_margin_factor)
86     history_cost.append(cost)
87
88     if epoch % 20 == 0 or epoch == max_epochs - 1:
89         print("Epoch is: {} and Cost is: {}".format(epoch, cost))
90
91
92
93
94     # Stoppage criterion
95     if prev_cost < cost:
96         if patience == 10:
97             return weights, history_cost
98         else:
99             patience += 1
100
101     else:
102         patience = 0
103         prev_cost = cost
104
105
106     return weights, history_cost
107
108     patience = 0
109     prev_cost = cost
110
111     return weights, history_cost
```





# Support vector machine

calculate **cost**

$$\text{cost}(w) = \frac{1}{2} ||w||^2 + C \left[ \frac{1}{N} \sum_i^n \max(0, 1 - y_i * (w * x_i + b)) \right]$$

calculate **cost-gradient**

$$\nabla_w \text{cost}(w) = \frac{1}{N} \sum_i^n \begin{cases} w & \text{if } \max(0, 1 - y_i * (w * x_i)) = 0 \\ w - C y_i x_i & \text{otherwise} \end{cases}$$

minimize gradient using **stochastic gradient decent**

train model

# *Our Algorithm*

Segmentation

Dataframe with normalized  
features and ground truths

SVM

Segmented images


Evaluation


Dice score



# Evaluation using the dice score

```
def dice_score(pred, gt):  
    """  
    This function calculates the similarity between two arrays.  
    :param pred: an array of predicted labels  
    :param gt: the ground truth of this array  
    :return: a value between 0 and 1, describing the similarity between those arrays. 1 is the dice score of similar arrays.  
    """  
    dice = np.sum(pred[gt == pred]) * 2.0 / (np.sum(gt) + np.sum(pred))  
    print(dice)
```



$$\text{Dice} = \frac{2 * \text{Intersection}}{\text{Union} + \text{Intersection}} = F_1 = \frac{1}{\frac{1}{\text{Prec}} + \frac{1}{\text{Recall}}} = \frac{2 TP}{2 TP + FP + FN}$$




Unittesting with synthetic masks  
and  
manually created arrays

# Results: N2DH-GOW1


original image:  
t21.tif




ground truth:  
man\_seg21.tif



segmented image:  
t21\_seg.tif  
dice score: 0.67




# Results: N2DH-GOW1



Features	Dice Score (DS)
no filters (NF)	good DS, but highest variance
all filters	worse than NF, but lower variance
Gauss, Otsu, PCA	worse than NF, but lower variance
Watershed	best DS

- average dice score with all filters: 0.71


# Results: N2DL-HeLa



Features	Dice Score (DS)
no filters (NF), Watershed	no segmentation
all filters	best DS
Otsu	biggest DS improvement
Gauss, PCA	not as relevant

- average Dice score with all filters: 0.63

# Results: NIH3T3



Features	Dice Score (DS)
no filters (NF)	DS over 0.5
all filters, PCA	highest DS
Gauss, Watershed	similar to NF
Otsu	good DS, but high variance

- average Dice score with all filters: 0.65




# Optimal settings for our SVM

- Learning rate:  $1 \cdot 10^{-7}$
- Regularization parameter: 10.000
- Maximum epochs: 40


## ➤ Why?

- Different learning rates
- Regularization parameter: 10.000
- Maximum epochs: 40


Learning rate:  $1 \cdot 10^{-5}$



Learning rate:  $1 \cdot 10^{-7}$



Learning rate:  $1 \cdot 10^{-9}$




# Optimal settings for our SVM

- Learning rate:  $1 \cdot 10^{-7}$
- Regularization parameter: 10.000
- Maximum epochs: 40


## ➤ Why?

- Learning rate:  $1 \cdot 10^{-7}$
- Different regularization parameters
- Maximum epochs: 40


Regularization parameter: 10.000.000



Regularization parameter: 10.000



Regularization parameter: 100




# Optimal settings for our SVM

- Learning rate:  $1 \cdot 10^{-7}$
- Regularization parameter: 10.000
- Maximum epochs: 40


## ➤ Why?

- Learning rate:  $1 \cdot 10^{-7}$
- Regularization parameter: 10.000
- Different maximum epochs

Maximum epochs: 100



Maximum epochs: 40



# *What we've learned*

- pre-processing methods depend on specific challenges of the data
- SVM with a Linear kernel is easier to code and takes less time to run
- SVM is a powerful tool for segmentation

But:

- needs a lot of computational power and thus has long runtime



# Outlook

SVMs include many possibilities for expansion

- Non-linear kernels for better segmentation results e.g. the RBF kernel
- include other features e.g. intensity values of the neighborhood
- test correlation of features to reduce runtime
- generate weights vector from changing starting points



**Thank you for  
listening!**

